

CSC108_HarmonyHealthcare_Data_Cleaning_and_Visualizations,_and_Lo 1

December 8, 2025

Mike and Spencer

0.0.1 Beyond the Emergency Department: Predictive Analytics for 6-Month Patient Admission in Harmony Healthcare

0.0.2 Problem Statement

Our goal is to build a predictive model that identifies whether a patient will experience at least one ED admission within the next six months. Based on instructor feedback, we removed columns directly tied to emergency department data, forcing the model to learn broader risk signals rather than simply reproducing ED activity. This notebook documents data cleaning, exploratory visualization, and an initial predictive model using forward feature selection with logistic regression.

```
[25]: # from google.colab import drive
# drive.mount('/content/drive')

import pandas as pd
df = pd.read_excel('HarmonyHealthcareOneWeek_9_2025.xlsx')

# We want to keep the 'ED Episode Admit Last-6-Mths' column and remove all other columns with ED in the name
ed_columns = [col for col in df.columns if 'ED' in col and col != 'ED Episode Admit Last-6-Mths']
df = df.drop(columns=ed_columns)

# Assuming NaN in 'ED Episode Admit Last-6-Mths' means no admission, fill with 0 early
target = 'ED Episode Admit Last-6-Mths'
df[target] = df[target].fillna(0)

# Now we can remove any columns that are just empty to shrink the data further
df = df.dropna(axis=1, how='all')
df.head()
```

[25]:

	Age	EHR	Sex	ED	Episode	Admit	Last-6-Mths	Most Recent	Encounter	Type	\
0	26	male					0.0			TeleBHTher	
1	14	female					0.0			Nutrition	
2	11	male					0.0			F/U	
3	49	male					0.0			BH Therapy	
4	39	female					0.0			Telemed	

	UDS	Qualifying	Encounter	Count	UDS	Homelessness	Status	\
0				21.0		Not Homeless		
1				4.0		Not Homeless		
2				5.0		Not Homeless		
3				48.0		Not Homeless		
4				4.0		Not Homeless		

	IP	Episode	Admit	Date	IP	Episode	Admit	Location	IP	Episode	Admit	Readmit	\
0				NaT				NaN				NaN	
1				NaT				NaN				NaN	
2				NaT				NaN				NaN	
3				NaT				NaN				NaN	
4				NaT				NaN				NaN	

	IP	Episode	Admit	Past-6-Mths	...	Urine	Creatinine	Date	\
0				NaN	...			NaT	
1				NaN	...			NaT	
2				NaN	...			NaT	
3				NaN	...			NaT	
4				NaN	...			NaT	

	Urine	Creatinine	Code	Urine	Creatinine	Result	Varicella	Titer	Date	\
0			NaN			NaN			NaT	
1			NaN			NaN			NaT	
2			NaN			NaN			NaT	
3			NaN			NaN			NaT	
4			NaN			NaN			NaT	

	Varicella	Titer	Code	Varicella	Titer	Result	Violence	Screening	Date	\
0			NaN			NaN		2025-02-10		
1			NaN			NaN			NaT	
2			NaN			NaN			NaT	
3			NaN			NaN		2025-08-26		
4			NaN			NaN		2025-07-26		

	Violence	Screening	Type	Vision	Screening	Date	Vision	Screening	Code
0	Domestic	Violence	PRAPARE			2024-11-08			99173
1				NaN		2025-01-02			99173
2				NaN		2024-11-09			99173
3	Domestic	Violence	PRAPARE			2024-09-25			99173

[5 rows x 583 columns]

```
[26]: target = 'ED Episode Admit Last-6-Mths' # Store our target var so we can use it later

# Check target distribution
print(df[target].value_counts())
print(df[target].value_counts(normalize=True) * 100)

# Check missing percentages
missing_pct = (df.isna().mean() * 100).sort_values(ascending=False)
print(missing_pct[missing_pct > 0])
```

```
ED Episode Admit Last-6-Mths
0.0      2442
1.0       170
2.0        39
3.0        10
4.0         4
5.0         4
6.0         1
34.0        1
7.0         1
Name: count, dtype: int64
ED Episode Admit Last-6-Mths
0.0      91.392216
1.0      6.362275
2.0      1.459581
3.0      0.374251
4.0      0.149701
5.0      0.149701
6.0      0.037425
34.0     0.037425
7.0      0.037425
Name: proportion, dtype: float64
Material Security Food-Insecurity-ICD10-Date      99.962575
Opioid Dependence Dx Name                          99.962575
Transportation ICD10-Insecurity-Code            99.962575
Transportation ICD10-Insecurity-Date            99.962575
Cerebral Palsy Date                            99.962575
...
SDOH Triggers                                0.898204
Patient Appointment No-Show Rate %           0.149701
Patient Appointment No-Show Count            0.149701
Patient Medicaid Risk Total Risk             0.074850
Patient Medicaid Risk Risk Gap              0.074850
```

Length: 561, dtype: float64

```
[27]: # Drop columns with >70% missing, except target
columns_to_drop = [c for c in df.columns if c != target and df[c].isna().mean() * 100 > 70]
print("Number of columns to drop:", len(columns_to_drop))
df = df.drop(columns=columns_to_drop)
df.head()
```

Number of columns to drop: 417

```
[27]:   Age EHR Sex ED Episode Admit Last-6-Mths Most Recent Encounter Type \
0    26 male           0.0             TeleBHTher
1    14 female         0.0             Nutrition
2    11 male           0.0             F/U
3    49 male           0.0             BH Therapy
4    39 female         0.0             Telemed

      UDS Qualifying Encounter Count UDS Homelessness Status Active Medications \
0          21.0            Not Homeless        13
1          4.0            Not Homeless        11
2          5.0            Not Homeless        23
3         48.0            Not Homeless        23
4          4.0            Not Homeless         4

      Alcohol Assessment Date Alcohol Assessment Code Alcohol Assessment Result \
0              2024-12-06          AUDIT-C           0.0
1              2025-09-17          AUDIT-C           0.0
2                  NaT            NaN             NaN
3              2024-06-05          AUDIT-C           3.0
4              2024-10-30          AUDIT-C           0.0

      ... Stress Response Transportation Trigger Transportation NonMed-Date \
0 ... Not at all           NaN            NaN             NaN
1 ... NaN                 NaN            NaN             NaN
2 ... NaN                 NaN            NaN             NaN
3 ... Not at all           N             8/26/2025 12:00:00 AM
4 ... A little bit         N             7/26/2025 12:00:00 AM

      Transportation NonMed-Reponse      UDS SDOH Triggers      UDS SDOH Tally \
0                   NaN          FINANCIAL STRAIN        1.0
1                   NaN          FINANCIAL STRAIN        1.0
2                   NaN          FINANCIAL STRAIN        1.0
3                   N FINANCIAL STRAIN FOOD       2.0
4                   N FINANCIAL STRAIN        1.0

      Violence Screening Date      Violence Screening Type Vision Screening Date \
0                2025-02-10 Domestic Violence PRAPARE           2024-11-08
```

```

1          NaT          NaN 2025-01-02
2          NaT          NaN 2024-11-09
3  2025-08-26 Domestic Violence PRAPARE 2024-09-25
4  2025-07-26 Domestic Violence PRAPARE 2025-06-24

    Vision Screening Code
0            99173
1            99173
2            99173
3            99173
4            99173

[5 rows x 166 columns]

```

```
[28]: # Impute missing values + encode categoricals

from sklearn.impute import SimpleImputer

target_col = 'ED Episode Admit Last-6-Mths' # Define target here too for consistency

# Separate numerical + categorical
numeric_cols = df.select_dtypes(include='number').columns.tolist()
# Exclude the target column from numeric imputation
if target_col in numeric_cols:
    numeric_cols.remove(target_col)

cat_cols = df.select_dtypes(include=['object', 'category']).columns.tolist()

# Impute numerics with median
num_imputer = SimpleImputer(strategy='median')
df[numeric_cols] = num_imputer.fit_transform(df[numeric_cols])

# Impute categoricals with most frequent
cat_imputer = SimpleImputer(strategy='most_frequent')
df[cat_cols] = cat_imputer.fit_transform(df[cat_cols])

# One-hot encode categoricals
df = pd.get_dummies(df, columns=cat_cols, drop_first=True)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2672 entries, 0 to 2671
Columns: 7049 entries, Age to Vision Screening Code_Z01.01
dtypes: bool(6970), datetime64[ns](40), float64(39)
memory usage: 19.4 MB
```


```
[29]: # Visualizations!
```


```
import matplotlib.pyplot as plt
import seaborn as sns

# 1. Target distribution
df[target].value_counts().plot(kind='bar')
plt.title("Target Distribution: 6-Month Admission")
plt.show()

# 2. Histogram of top numeric features
# df[numeric_cols].hist(figsize=(12,10))
# plt.tight_layout()
# plt.show()

# 3. Correlation heatmap
plt.figure(figsize=(12,10))
sns.heatmap(df[numeric_cols].corr(), cmap='coolwarm')
plt.title("Correlation Heatmap (Numeric Features)")
plt.show()
```





```
[30]: # Convert target to binary (admitted = 1 if >0)
y = (df[target] > 0).astype(int)
X = df.drop(columns=[target])
```

```
[31]: # Clean and prepare data

# Make a copy of the dataframe to avoid modifying the original `df` from the previous cell.
df_processed = df.copy()

# Target column name
target_col_name = "ED Episode Admit Last-6-Mths"
```

```

# Isolate the target variable BEFORE any potentially destructive
# transformations.
# Apply fillna(0) and >0 to ensure binary classification, and then convert to
# int.
y = (df_processed[target_col_name].fillna(0) > 0).astype(int)

# Drop the target column from the feature set X
X = df_processed.drop(columns=[target_col_name])

# Identify columns that are actual datetime objects in X.
# These are the columns from `df.info()` with dtypes `datetime64[ns]` from the
# previous step.
datetime_cols_in_X = X.select_dtypes(include=['datetime64[ns]']).columns.
    tolist()

# Convert identified datetime columns in X to numeric days since epoch.
# This avoids incorrectly converting other numeric columns.
for col in datetime_cols_in_X:
    X[col] = (X[col] - pd.Timestamp("1970-01-01")).dt.days

# Replace remaining missing numerical values in X with column medians.
# This imputation step is for the features (X) after date conversion.
X = X.fillna(X.median(numeric_only=True))

# Train/test split
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# Standardize numerical data
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```

[32]: # Logistic Regression + evaluation function
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score

def evaluate(feature_list):
    idxs = [X.columns.get_loc(f) for f in feature_list]

    model = LogisticRegression(
        max_iter=500,

```

```

        class_weight="balanced"
    )

model.fit(X_train_scaled[:, idxs], y_train)

preds = model.predict_proba(X_test_scaled[:, idxs])[:, 1]

return roc_auc_score(y_test, preds)

```

```
[33]: # Greedy forward feature selection

# Initialize lists: 'remaining' holds features not yet selected, 'selected' holds chosen features,
# and 'scores' stores the AUC for the selected feature set at each step.
remaining = list(X.columns)
selected = []
scores = []

# Perform 10 steps of forward feature selection.
# In each step, we find the single best feature to add to our 'selected' set.
for step in range(10):
    best_feature = None
    best_auc = -1 # Initialize with a low AUC score to ensure the first valid AUC is higher

    # Iterate through all features not yet selected to find the one that maximizes AUC when added.
    for feat in remaining:
        # Create a temporary list of features that includes currently selected features plus one 'candidate' feature.
        try_features = selected + [feat]
        # Evaluate the performance (AUC) of the model using this candidate set of features.
        auc = evaluate(try_features)

        # If this candidate set yields a better AUC than the current best, update best_auc and best_feature.
        if auc > best_auc:
            best_auc = auc
            best_feature = feat

    # Add the best performing feature from this step to the 'selected' list.
    selected.append(best_feature)
    # Record the AUC achieved with this new set of selected features.
    scores.append(best_auc)
    # Remove the selected feature from the 'remaining' list so it's not considered again.

```

```

remaining.remove(best_feature)

# Print the result for the current step.
# print(f"Step {step+1}: Selected {best_feature} - AUC {best_auc:.4f}")

# After all steps are complete, print the final list of top 10 selected
# features and their corresponding AUC scores.
print("\nTop 10 selected features:")
for i, (feat, auc) in enumerate(zip(selected, scores), 1):
    print(f"{i}. {feat} - AUC {auc:.4f}")

```

Top 10 selected features:

1. Patient Medicaid Risk Total Risk - AUC 0.6063
2. Insurance Primary Payer_HealthFirst MCD - AUC 0.6410
3. Most Recent Encounter Type_EOB - AUC 0.6611
4. Domestic Violence Reponse_I choose not to answer this question - AUC 0.6789
5. Sexually Active Code_43305-2 - AUC 0.6964
6. HCV Antibody Test Date - AUC 0.7141
7. Flu Date - AUC 0.7279
8. Last Well Care Visit Code_Z02.1 - AUC 0.7394
9. AST Result - AUC 0.7505
10. FPL Date_11/19/2024 12:00:00 AM - AUC 0.7591

0.0.3 Discussion of Early Results and Moving Forward

- The greedy selection approach identified a ranked set of features contributing incremental predictive value.
- The AUC values supply an interpretable measure of classification.
- Further improvements planned include:
 1. Cross validate the model with testing data
 2. Plot the AUC for more than 10 features
 3. Check correlation of selected features to make sure we are not selecting highly correlated features
 4. Attempt to compare to Lasso if enough time

The project was a collaborative effort between Mike and Spencer with both of us working on the data cleaning step and Mike taking charge of the algorithm and Spencer doing the write up and github steps. What remains is a little of a first come first serve effort with communication being key to make sure we don't step on each others toes.

Github link: <https://github.com/99x5zbrvgj-droid/CSC108HHCFinal>